

Demystifying the Nuts & Bolts of Kubernetes Architecture

ReplicaSet101



How can you ensure that there are 3 Pods instances which are always available and running at one point in time?



What is ReplicaSet all about?

Maintain a stable set of replica Pods running at any given time

- Ensures that a specified number of Pods are running at any time
 - a. If there are excess Pods, they get killed and vice versa
 - b. New Pods are launched when they get failed, get deleted and terminated
- ReplicaSet & Pods are associated with “labels”


Replication Controller Vs ReplicaSet

- ReplicaSet is the next generation of Replication Controller
- Both serve the same purpose



Labels & Selectors

When Pods are scaled, how are these Pods Managed at such large scale?



Equality-based Selectors

Operators:

= and ==

Examples:

environment = production
tier = frontend

Commandline:

```
$kubectl get pods -l environment=production
```

In Manifest:

```
..  
selector:  
  environment: production  
  tier: frontend  
..
```

Supports: Services, Replication Controller

Set-based Selectors

Operators:

in notin exists

Examples:

environment in (production, qa)
tier notin(frontend, backend)

Commandline:

```
$kubectl get pods -l `environment in(production)
```

In Manifest:

```
..  
selector:  
  matchExpressions:  
    - {key:environment,operator:in,values:[prod,qa]}  
    - {key:tier,operator:Notin,values:[frontend,backend]}  
..
```

Supports: Job, Deployment, ReplicaSet, DaemonSet

```
...
selector:
  app: nginx
  tier: frontend
...
```

=

```
...
selector:
  matchLabels:
    app: nginx
    tier: frontend
...
```

Supports on Older Resources such as:

- ReplicationControllers,
- Services

Supports on newer resources such as:

- ReplicaSets
- Deployments
- Jobs
- DaemonSet

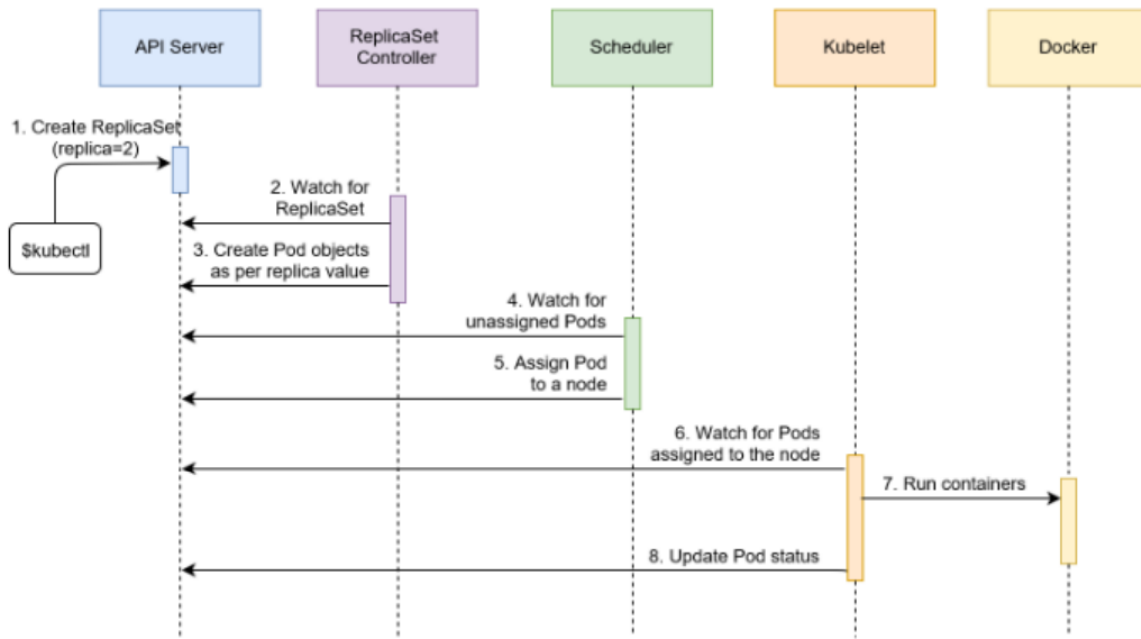
ReplicaSet Examples:

- Manifest file
- Deploy app using RS
- Display and validate RS
- Test – Node Fails
- Test – Scale Up
- Test – Scale Down

ReplicaSet Manifest File

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-rs
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx-app
  template:
    metadata:
      name: nginx-pod
      labels:
        app: nginx-app
        tier: frontend
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

A Typical Replicaset Workflow



Credits" Viktor Farcic

Creating Nginx-rs Pods

\$kubectl create -f nginx-rs.yaml

```
[node1 lab02-creating-replicaset]$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-pod	1/1	Running	0	36m
nginx-rs-jl266	1/1	Running	0	62s
nginx-rs-jq74j	1/1	Running	0	62s

```
[node1 lab02-creating-replicaset]$ kubectl get po -l tier=frontend
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-rs-jl266	1/1	Running	0	2m52s
nginx-rs-jq74j	1/1	Running	0	2m52s

```
[node1 lab02-creating-replicaset]$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-rs	2	2	1	12m

```
[node1 lab02-creating-replicaset]$ kubectl get rs -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
nginx-rs	2	2	1	12m	nginx	nginx	app=nginx-app

Checking the state of ReplicaSet

```
[node1 lab02-creating-replicaset]$ kubectl describe rs
Name:          nginx-rs
Namespace:     default
Selector:      app=nginx-app
Labels:        <none>
Annotations:   <none>
Replicas:      2 current / 2 desired
Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=nginx-app
           tier=frontend
  Containers:
    nginx:
      Image:          nginx
      Port:           80/TCP
      Host Port:      0/TCP
      Environment:    <none>
      Mounts:         <none>
      Volumes:        <none>
Events:
  Type     Reason             Age   From                  Message
  ----     -
  Normal   SuccessfulCreate   14m   replicaset-controller Created pod: nginx-rs-jq74j
  Normal   SuccessfulCreate   14m   replicaset-controller Created pod: nginx-rs-jl266
```

Scaling the Nginx Service

```
[node1 lab02-creating-replicaset]$ kubectl scale rs nginx-rs --replicas=5  
replicaset.extensions/nginx-rs scaled
```

Demo

- [Introductory Slides](#)
- [Creating Your First ReplicaSet - 4 Pods serving Nginx](#)
- [Removing a Pod from ReplicaSet](#)
- [Scaling & Autoscaling a ReplicaSet](#)
- [Best Practices](#)
- [Deleting ReplicaSets](#)

Reference

- <https://kubelabs.collabnix.com>
- <https://kubetools.collabnix.com>

Thank You