



Deployment 101

Kubernetes - Beginners | Intermediate | Advanced

[View on GitHub](#) [Join Slack](#)

Deployment 101

We looked at ReplicaSets earlier. However, ReplicaSet have one major drawback: once you select the pods that are managed by a ReplicaSet, you cannot change their pod templates.

For example, if you are using a ReplicaSet to deploy four pods with NodeJS running and you want to change the NodeJS image to a newer version, you need to delete the ReplicaSet and recreate it. Restarting the pods causes downtime till the images are available and the pods are running again.

A Deployment resource uses a ReplicaSet to manage the pods. However, it handles updating them in a controlled way. Let's dig deeper into Deployment Controllers and patterns.

Creating Your First Deployment

The following Deployment definition deploys four pods with nginx as their hosted application:

```
git clone https://github.com/collabnix/dockerlabs
cd dockerlabs/kubernetes/workshop/Deployment101
kubectl create -f nginx-dep.yaml
deployment.apps/nginx-deployment created
```

Checking the list of application deployment

To list your deployments use the get deployments command:

```
$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	2/2	2	2	63s

```
[node1 Deployment101]$ kubectl describe deploy
```

Name: nginx-deployment
Namespace: default
CreationTimestamp: Mon, 30 Dec 2019 07:10:33 +0000
Labels: <none>
Annotations: deployment.kubernetes.io/revision: 1
Selector: app=nginx
Replicas: 2 desired | 2 updated | 2 total | 0 available | 2 un
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
 Labels: app=nginx
 Containers:
 nginx:
 Image: nginx:1.7.9
 Port: 80/TCP
 Host Port: 0/TCP
 Environment: <none>
 Mounts: <none>
 Volumes: <none>
Conditions:
 Type Status Reason

```

----
Available      False    MinimumReplicasUnavailable
Progressing    True     ReplicaSetUpdated
OldReplicaSets: <none>
NewReplicaSet:  nginx-deployment-6dd86d77d (2/2 replicas created)
Events:
  Type          Reason              Age   From                      Message
  ----          -
  Normal        ScalingReplicaSet   90s   deployment-controller     Scaled up replica

```

We should have 1 Pod. If not, run the command again. This shows:

```

The DESIRED state is showing the configured number of replicas
The CURRENT state show how many replicas are running now
The UP-TO-DATE is the number of replicas that were updated to match the desi
The AVAILABLE state shows how many replicas are actually AVAILABLE to the us

```

```

[node1 Deployment101]$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    2/2     2             2           2m57s

```

```

[node1 Deployment101]$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6dd86d77d-84fwp    1/1     Running   0           3m44s
nginx-deployment-6dd86d77d-xnrqp    1/1     Running   0           3m44s

```

Step #2. Scale up/down application deployment

Now let's scale the Deployment to 4 replicas. We are going to use the `kubectl scale` command, followed by the deployment type, name and desired number of instances:

```

Biradars-MacBook-Air-4:~ sangam$ kubectl scale deployments/nginx-deployment
deployment.extensions/nginx-deployment scaled

```

The change was applied, and we have 4 instances of the application available. Next, let's

check if the number of Pods changed:

Now There should be 4 pods running in the cluster

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    4/4     4            4           4m
```

There are 4 Pods now, with different IP addresses. The change was registered in the Deployment events log. To check that, use the describe command:

```
$ kubectl describe deployments/nginx-deployment
Name:                nginx-deployment
Namespace:           default
CreationTimestamp:    Sat, 30 Nov 2019 20:04:34 +0530
Labels:              <none>
Annotations:         deployment.kubernetes.io/revision: 1
Selector:             app=nginx
Replicas:            4 desired | 4 updated | 4 total | 4 available | 0 un
StrategyType:        RollingUpdate
MinReadySeconds:     0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:        nginx:1.7.9
      Port:         80/TCP
      Host Port:    0/TCP
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>
Conditions:
  Type           Status  Reason
  ----           -
  Progressing    True    NewReplicaSetAvailable
  Available      True    MinimumReplicasAvailable
OldReplicaSets:  <none>
```

```
NewReplicaSet:  nginx-deployment-6dd86d77d (4/4 replicas created)
Events:
  Type      Reason              Age   From                    Message
  ----      -
  Normal    ScalingReplicaSet   6m12s deployment-controller   Scaled up replica
  Normal    ScalingReplicaSet   3m6s  deployment-controller   Scaled up replica
Biradars-MacBook-Air-4:~ sangam$
```

```
$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP
nginx-deployment-6dd86d77d-b4v7k    1/1     Running   0           4m32s  10.1
nginx-deployment-6dd86d77d-bnc5m    1/1     Running   0           4m32s  10.1
nginx-deployment-6dd86d77d-bs6jr    1/1     Running   0           86s    10.1
nginx-deployment-6dd86d77d-wbdzv    1/1     Running   0           86s    10.1
Biradars-MacBook-Air-4:~ sangam$
```

You can also view in the output of this command that there are 4 replicas now.

Scaling the service to 2 Replicas

To scale down the Service to 2 replicas, run again the scale command:

```
$ kubectl scale deployments/nginx-deployment --replicas=2
deployment.extensions/nginx-deployment scaled
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    2/2     2             2           7m23s
Biradars-MacBook-Air-4:~ sangam$
```

Step #3. Perform rolling updates to application deployment

So far, everything our Deployment did is no different than a typical ReplicaSet. The real power of a Deployment lies in its ability to update the pod templates without causing application outage.

Let's say that you have finished testing the nginx 1.7.9 , and you are ready to use it in

production. The current pods are using the older nginx version . The following command changes the deployment pod template to use the new image:

To update the image of the application to new version, use the set image command, followed by the deployment name and the new image version:

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    2/2     2             2           7m23s
$ kubectl describe pods
Name:                nginx-deployment-6dd86d77d-b4v7k
Namespace:           default
Priority:             0
PriorityClassName:    <none>
Node:                docker-desktop/192.168.65.3
Start Time:          Sat, 30 Nov 2019 20:04:34 +0530
Labels:              app=nginx
                    pod-template-hash=6dd86d77d
Annotations:         <none>
Status:              Running
IP:                  10.1.0.237
Controlled By:       ReplicaSet/nginx-deployment-6dd86d77d
Containers:
  nginx:
    Container ID:     docker://2c739cf9fe4dac53a4cc5c6097207da0c5edc2183f1f36f
    Image:             nginx:1.7.9
    Image ID:          docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fc
    Port:              80/TCP
    Host Port:         0/TCP
    State:             Running
      Started:         Sat, 30 Nov 2019 20:05:28 +0530
    Ready:             True
    Restart Count:     0
    Environment:       <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-ds5tg
Conditions:
  Type              Status
  Initialized       True
  Ready             True
```

```

ContainersReady    True
PodScheduled       True
Volumes:
  default-token-ds5tg:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-ds5tg
    Optional:      false
QoS Class:         BestEffort
Node-Selectors:    <none>
Tolerations:       node.kubernetes.io/not-ready:NoExecute for 300s
                   node.kubernetes.io/unreachable:NoExecute for 300s

Events:
  Type        Reason          Age   From                      Message
  ----        -
  Normal      Scheduled       10m   default-scheduler        Successfully assigned d
  Normal      Pulling         10m   kubelet, docker-desktop  Pulling image "nginx:1.
  Normal      Pulled          9m17s kubelet, docker-desktop  Successfully pulled ima
  Normal      Created         9m17s kubelet, docker-desktop  Created container nginx
  Normal      Started         9m17s kubelet, docker-desktop  Started container nginx

Name:              nginx-deployment-6dd86d77d-bnc5m
Namespace:         default
Priority:           0
PriorityClassName: <none>
Node:              docker-desktop/192.168.65.3
Start Time:        Sat, 30 Nov 2019 20:04:34 +0530
Labels:            app=nginx
                   pod-template-hash=6dd86d77d
Annotations:       <none>
Status:            Running
IP:                10.1.0.236
Controlled By:     ReplicaSet/nginx-deployment-6dd86d77d
Containers:
  nginx:
    Container ID:   docker://12ab35cbf4fdf78997b106b5eb27135f2fc37c890e723fe
    Image:          nginx:1.7.9
    Image ID:       docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fc
    Port:           80/TCP
    Host Port:      0/TCP
    State:          Running

```

```

    Started:      Sat, 30 Nov 2019 20:05:23 +0530
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-ds5tg
Conditions:
  Type            Status
  Initialized      True
  Ready            True
  ContainersReady  True
  PodScheduled     True
Volumes:
  default-token-ds5tg:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-ds5tg
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From                    Message
  ----    -
  Normal  Scheduled   10m   default-scheduler      Successfully assigned d
  Normal  Pulling     10m   kubelet, docker-desktop Pulling image "nginx:1.
  Normal  Pulled      9m22s kubelet, docker-desktop Successfully pulled ima
  Normal  Created     9m22s kubelet, docker-desktop Created container nginx
  Normal  Started     9m22s kubelet, docker-desktop Started container nginx

```

The command notified the Deployment to use a different image for your app and initiated a rolling update. Check the status of the new Pods, and view the old one terminating with the `get pods` command:

```

Biradars-MacBook-Air-4:~ sangam$ kubectl set image deployments/nginx-deploy
deployment.extensions/nginx-deployment image updated

```

Checking description of pod again


```
$ kubectl describe pods
Name:          nginx-deployment-6dd86d77d-b4v7k
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          docker-desktop/192.168.65.3
Start Time:    Sat, 30 Nov 2019 20:04:34 +0530
Labels:        app=nginx
               pod-template-hash=6dd86d77d
Annotations:   <none>
Status:        Running
IP:            10.1.0.237
Controlled By: ReplicaSet/nginx-deployment-6dd86d77d
Containers:
  nginx:
    Container ID:  docker://2c739cf9fe4dac53a4cc5c6097207da0c5edc2183f1f36f
    Image:         nginx:1.7.9
    Image ID:      docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fc
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Sat, 30 Nov 2019 20:05:28 +0530
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-ds5tg
Conditions:
  Type          Status
  Initialized    True
  Ready         True
  ContainersReady True
  PodScheduled   True
Volumes:
  default-token-ds5tg:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-ds5tg
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
```

Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
node.kubernetes.io/unreachable:NoExecute for 300s

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	Scheduled	16m	default-scheduler	Successfully assigned de
Normal	Pulling	16m	kubelet, docker-desktop	Pulling image "nginx:1.7
Normal	Pulled	15m	kubelet, docker-desktop	Successfully pulled imag
Normal	Created	15m	kubelet, docker-desktop	Created container nginx
Normal	Started	15m	kubelet, docker-desktop	Started container nginx

Name: nginx-deployment-6dd86d77d-bnc5m

Namespace: default

Priority: 0

PriorityClassName: <none>

Node: docker-desktop/192.168.65.3

Start Time: Sat, 30 Nov 2019 20:04:34 +0530

Labels: app=nginx
pod-template-hash=6dd86d77d

Annotations: <none>

Status: Running

IP: 10.1.0.236

Controlled By: ReplicaSet/nginx-deployment-6dd86d77d

Containers:

nginx:

Container ID:	docker://12ab35cbf4fdf78997b106b5eb27135f2fc37c890e723fe
Image:	nginx:1.7.9
Image ID:	docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fc
Port:	80/TCP
Host Port:	0/TCP
State:	Running
Started:	Sat, 30 Nov 2019 20:05:23 +0530
Ready:	True
Restart Count:	0
Environment:	<none>

Mounts:

/var/run/secrets/kubernetes.io/serviceaccount from default-token-ds5tg

Conditions:

Type	Status
Initialized	True

```

    Ready                True
    ContainersReady      True
    PodScheduled         True
Volumes:
  default-token-ds5tg:
    Type:                Secret (a volume populated by a Secret)
    SecretName:          default-token-ds5tg
    Optional:            false
QoS Class:              BestEffort
Node-Selectors:         <none>
Tolerations:            node.kubernetes.io/not-ready:NoExecute for 300s
                       node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   16m   default-scheduler   Successfully assigned de
  Normal  Pulling     16m   kubelet, docker-desktop   Pulling image "nginx:1.7
  Normal  Pulled      15m   kubelet, docker-desktop   Successfully pulled imag
  Normal  Created     15m   kubelet, docker-desktop   Created container nginx
  Normal  Started     15m   kubelet, docker-desktop   Started container nginx

Name:          nginx-deployment-784b7cc96d-kxc68
Namespace:    default
Priority:      0
PriorityClassName: <none>
Node:         docker-desktop/192.168.65.3
Start Time:   Sat, 30 Nov 2019 20:20:04 +0530
Labels:       app=nginx
              pod-template-hash=784b7cc96d
Annotations:  <none>
Status:       Pending
IP:
Controlled By: ReplicaSet/nginx-deployment-784b7cc96d
Containers:
  nginx:
    Container ID:
    Image:        nginx:1.9.1
    Image ID:
    Port:         80/TCP
    Host Port:    0/TCP

```

```

State:      Waiting
Reason:     ContainerCreating
Ready:      False
Restart Count: 0
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-ds5tg
Conditions:
  Type              Status
  Initialized       True
  Ready             False
  ContainersReady   False
  PodScheduled      True
Volumes:
  default-token-ds5tg:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-ds5tg
    Optional:       false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From              Message
  ----    -
  Normal  Scheduled   36s   default-scheduler Successfully assigned de
  Normal  Pulling     35s   kubelet, docker-desktop Pulling image "nginx:1.9
Biradars-MacBook-Air-4:~ sangam$

```

Step #4. Rollback updates to application deployment

The rollout command reverted the deployment to the previous known state. Updates are versioned and you can revert to any previously know state of a Deployment. List again the Pods:

```

$ kubectl rollout undo deployments/nginx-deployment
deployment.extensions/nginx-deployment rolled back

$ kubectl rollout status deployments/nginx-deployment

```

```
deployment "nginx-deployment" successfully rolled out
```

After the rollout succeeds, you may want to get the Deployment.

The output shows the update progress until all the pods use the new container image.

The algorithm that Kubernetes Deployments use when deciding how to roll updates is to keep at least 25% of the pods running. Accordingly, it doesn't kill old pods unless a sufficient number of new ones are up. In the same sense, it does not create new pods until enough pods are no longer running. Through this algorithm, the application is always available during updates.

You can use the following command to determine the update strategy that the Deployment is using:

```
Biradars-MacBook-Air-4:~ sangam$ kubectl describe deployments | grep Strateg
StrategyType:          RollingUpdate
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Biradars-MacBook-Air-4:~ sangam$
```

Step #5. Cleanup

Finally you can clean up the resources you created in your cluster:

```
kubectl delete service nginx-deployment
kubectl delete deployment nginx-deployment
```

Contributors

Sangam Biradar

Reviewers

Ajeet Singh Raina

[Next »](#)


Join KubeDaily

9 Members Online

Support

MEMBERS ONLINE

 Aman Manapure


 h3ll_boY


 MEE6

 Nitinkashyap


Apex Legends

 ojaswa

 Parmeshwar

 prasad

 trimankaur

 vikas027

Free voice chat from Discord

Connect

[Tweets by collabnix](#)

kubelabs is maintained by [collabnix](#).

This page was generated by [GitHub Pages](#).