



Star	898	Fork	152	Watch	52	Follow @collabnix	0
------	-----	------	-----	-------	----	-------------------	---

What is node affinity ?

Kubernetes - Beginners | Intermediate | Advanced

[View on GitHub](#) [Join Slack](#)

What is node affinity ?

- In simple words this allows you to tell Kubernetes to schedule pods only to specific subsets of nodes.
- The initial node affinity mechanism in early versions of Kubernetes was the `nodeSelector` field in the pod specification. The node had to include all the labels specified in that field to be eligible to become the target for the pod.

nodeSelector

Steps

```
git clone https://github.com/collabnix/dockerlabs
cd dockerlabs/kubernetes/workshop/Scheduler101/
kubectl label nodes node2 mynode=worker-1
kubectl apply -f pod-nginx.yaml
```

- We have label on the node with node name, in this case i have given node2 as

mynode=worker-1 label.

Viewing Your Pods

```
kubectl get pods --output=wide
```

```
[node1 Scheduler101]$ kubectl describe po nginx
Name:          nginx
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          node2/192.168.0.17
Start Time:    Mon, 30 Dec 2019 16:40:53 +0000
Labels:        env=test
Annotations:   kubectl.kubernetes.io/last-applied-configuration:
                {"apiVersion":"v1","kind":"Pod","metadata":{"annotation
Status:        Pending
IP:
Containers:
  nginx:
    Container ID:
    Image:        nginx
    Image ID:
    Port:         <none>
    Host Port:    <none>
    State:        Waiting
      Reason:     ContainerCreating
    Ready:        False
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-qpgxo
Conditions:
  Type           Status
  Initialized     True
  Ready           False
  ContainersReady False
  PodScheduled    True
```

```

Volumes:
  default-token-qpgxq:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-qpgxq
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  mynode=worker-1
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type       Reason          Age   From                  Message
  ----       -
Normal      Scheduled       7s    default-scheduler     Successfully assigned default/
Normal      Pulling         3s    kubelet, node2        Pulling image "nginx"
[node1 Scheduler101]$

```

- You can check in above output Node-Selectors: mynode=worker-1

Deleting the Pod

```

kubectl delete -f pod-nginx.yaml
pod "nginx" deleted

```

Node affinity

- Node affinity is conceptually similar to nodeSelector – it allows you to constrain which nodes your pod is eligible to be scheduled on, based on labels on the node.
- There are currently two types of node affinity.
 1. requiredDuringSchedulingIgnoredDuringExecution (Preferred during scheduling, ignored during execution; we are also known as “hard” requirements)
 2. preferredDuringSchedulingIgnoredDuringExecution (Required during scheduling, ignored during execution; we are also known as “soft”

requirements)

Steps

```
git clone https://github.com/collabnix/dockerlabs
cd dockerlabs/kubernetes/workshop/Scheduler101/
kubectl label nodes node2 mynode=worker-1
kubectl label nodes node3 mynode=worker-3
kubectl apply -f pod-with-node-affinity.yaml
```

Viewing Your Pods

```
kubectl get pods --output=wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
with-node-affinity	1/1	Running	0	9m46s	10.44.0.1	kube-s

```
[node1 Scheduler101]$ kubectl describe po
Name:                with-node-affinity
Namespace:           default
Priority:             0
PriorityClassName:    <none>
Node:                node3/192.168.0.16
Start Time:          Mon, 30 Dec 2019 19:28:33 +0000
Labels:              <none>
Annotations:         kubectl.kubernetes.io/last-applied-configuration:
                      {"apiVersion":"v1","kind":"Pod","metadata":{"annotation
Status:              Pending
IP:
Containers:
  nginx:
    Container ID:
    Image:          nginx
    Image ID:
    Port:           <none>
```

```

Host Port:      <none>
State:         Waiting
  Reason:      ContainerCreating
Ready:         False
Restart Count: 0
Environment:   <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-qpgxq
Conditions:
  Type              Status
  Initialized        True
  Ready              False
  ContainersReady    False
  PodScheduled       True
Volumes:
  default-token-qpgxq:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-qpgxq
    Optional:      false
QoS Class:        BestEffort
Node-Selectors:   <none>
Tolerations:      node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   26s   default-scheduler  Successfully assigned default/
  Normal  Pulling     22s   kubelet, node3     Pulling image "nginx"
  Normal  Pulled      20s   kubelet, node3     Successfully pulled image "ngi
  Normal  Created     2s    kubelet, node3     Created container nginx
  Normal  Started     0s    kubelet, node3     Started container nginx

```

Step Cleanup

Finally you can clean up the resources you created in your cluster:

```
kubectl delete -f pod-with-node-affinity.yaml
```

[Next»](#)


Join KubeDaily

9 Members Online

Support

MEMBERS ONLINE


 Aman Manapure


 h3ll_boY


 MEE6


 Nitinkashyap

Apex Legends

 ojaswa

 Parmeshwar

 prasad

 trimankaur

 vikas027

Free voice chat from Discord

Connect

[Tweets by collabnix](#)

kubelabs is maintained by **collabnix**.

This page was generated by [GitHub Pages](#).