



Star	898	Fork	152	Watch	52	Follow @collabnix	0
------	-----	------	-----	-------	----	-------------------	---

Anti-Node Affinity ?

Kubernetes - Beginners | Intermediate | Advanced

[View on GitHub](#) [Join Slack](#)

Anti-Node Affinity ?

- Some scenarios require that you don't use one or more nodes except for particular pods. Think of the nodes that host your monitoring application.
- Those nodes shouldn't have many resources due to the nature of their role. Thus, if other pods than those which have the monitoring app are scheduled to those nodes, they hurt monitoring and also degrades the application they are hosting.
- In such a case, you need to use node anti-affinity to keep pods away from a set of nodes.

Steps

```
git clone https://github.com/collabnix/dockerlabs
cd dockerlabs/kubernetes/workshop/Scheduler101/
kubectl label nodes node2 mynode=worker-1
kubectl label nodes node3 mynode=worker-3
kubectl apply -f pod-anti-node-affinity.yaml
```

Viewing Your Pods

```
[node1 Scheduler101]$ kubectl get pods --output=wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED N
nginx	1/1	Running	0	2m37s	10.44.0.1	node2	<none>

Get nodes label detail

```
[node1 Scheduler101]$ kubectl get nodes --show-labels | grep mynode
```

node2	Ready	<none>	166m	v1.14.9	beta.kubernetes.io/arch=amd64,bet
node3	Ready	<none>	165m	v1.14.9	beta.kubernetes.io/arch=amd64,bet

Get pod describe

```
[node1 Scheduler101]$ kubectl describe pods nginx
```

```
Name:          nginx
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          node2/192.168.0.17
Start Time:    Mon, 30 Dec 2019 19:02:46 +0000
Labels:        <none>
Annotations:   kubectl.kubernetes.io/last-applied-configuration:
                {"apiVersion":"v1","kind":"Pod","metadata":{"annotation
Status:        Running
IP:            10.44.0.1
Containers:
  nginx:
    Container ID:  docker://2bdc20d79c360e1cd857eeb9bbb9424c726b2133e78f25b
    Image:         nginx
    Image ID:      docker-pullable://nginx@sha256:b2d89d0a210398b4d1120b3e3
    Port:          <none>
    Host Port:     <none>
```

```

State:      Running
  Started:   Mon, 30 Dec 2019 19:03:07 +0000
Ready:      True
Restart Count: 0
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-qpgxq
Conditions:
  Type              Status
  Initialized        True
  Ready              True
  ContainersReady    True
  PodScheduled       True
Volumes:
  default-token-qpgxq:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-qpgxq
    Optional:       false
QoS Class:         BestEffort
Node-Selectors:    <none>
Tolerations:       node.kubernetes.io/not-ready:NoExecute for 300s
                   node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type       Reason      Age   From          Message
  ----       -
  Normal     Scheduled   60s   default-scheduler   Successfully assigned default/
  Normal     Pulling     56s   kubelet, node2      Pulling image "nginx"
  Normal     Pulled      54s   kubelet, node2      Successfully pulled image "ngi
  Normal     Created     40s   kubelet, node2      Created container nginx
  Normal     Started     39s   kubelet, node2      Started container nginx

```

- Adding another key to the matchExpressions with the operator NotIn will avoid scheduling the nginx pods on any node labelled worker-1.

Step Cleanup

Finally you can clean up the resources you created in your cluster:

```
kubectl delete -f pod-anti-node-affinity.yaml
```


[Next »](#)


Join KubeDaily

9 Members Online


Support

MEMBERS ONLINE


 Aman Manapure

 h3ll_b0y


 MEE6

 Nitinkashyap


Apex Legends

 ojaswa

 Parmeshwar

 prasad

 trimankaur

 vikas027

Free voice chat from Discord

Connect

[Tweets by collabnix](#)

kubelabs is maintained by collabnix.

This page was generated by [GitHub Pages](#).